

ScoreStream Backend Documentation

1) Executive Summary

ScoreStream is a Node.js backend for a live sports score and commentary platform. It exposes a REST API for match and commentary CRUD-style operations, persists data in PostgreSQL through Drizzle ORM, and pushes real-time events over WebSockets. The backend is protected with Arcjet security policies (bot detection, shielding, and rate limiting) across both HTTP requests and WebSocket upgrades.

At a high level:

- Express + HTTP server powers REST endpoints.
- WebSocket server (ws) is attached to the same HTTP server for live fan-out.
- PostgreSQL + Drizzle ORM handles persistence and query composition.
- Zod schemas enforce request-level validation.
- Arcjet enforces request protection and abuse controls.

2) Architecture Overview

2.1 System Context



2.2 Runtime Composition

- A single Node process creates:
 - express() app
 - http.createServer(app)
 - WebSocket server attached via `attachWebSocketServer(server)`
- REST routes publish domain events by calling broadcast functions saved in `app.locals`.

This design keeps the API and real-time transport tightly integrated and avoids cross-process event buses for simple deployments.

3) Technology Stack

3.1 Core Platform

- **Node.js (ESM mode)**: package is configured with "type": "module".
- **Express 5**: HTTP routing and middleware.
- **ws**: WebSocket transport and subscription channel implementation.

3.2 Data Layer

- **PostgreSQL (pg)**: primary relational database.
- **Drizzle ORM**: schema-first definitions + query builder for Node/Postgres.
- **Drizzle Kit**: migration generation, migration execution, and studio tooling.

3.3 Validation & Security

- **Zod**: strict input validation for params, query, and request bodies.
- **Arcjet**: abuse prevention for both HTTP and WebSocket upgrade paths.

3.4 Utility / Configuration

- `dotenv`: environment variable loading.

4) Project Structure and Responsibilities

```

src/
  index.js          # App bootstrap + server wiring
  arcjet.js         # Arcjet policy definitions + HTTP middleware
  db/
    db.js           # Postgres pool + Drizzle instance
    schema.js        # Match/commentary schema + enum
  routes/
    matches.js      # /matches REST handlers
    commentary.js   # /matches/:id/commentary REST handlers
  validation/
    matches.js      # Zod schemas for match endpoints
    commentary.js   # Zod schemas for commentary endpoints
  utils/
    match.status.js # Match lifecycle status derivation helpers
  ws/
    server.js       # WebSocket server + subscriptions + broadcasting

drizzle/
  *.sql            # generated SQL migrations

```

5) Request Lifecycle and Flow

5.1 HTTP Path

1. Request enters Express app.
2. `express.json()` parses JSON payload.
3. Arcjet HTTP middleware runs (`securityMiddleware()`):
 - o denylist/shield checks
 - o bot handling
 - o sliding window rate limiting
4. Route-level Zod validation executes.
5. Route handler performs DB operation through Drizzle.
6. Successful mutation may trigger real-time broadcast via `app.locals` callbacks.
7. JSON response returned to client.

5.2 WebSocket Path

1. Client upgrades to `/ws`.
2. Upgrade request passes Arcjet WS protection.
3. On connection:
 - o client receives { `type: "welcome"` }
 - o heartbeat (ping/pong) liveness tracking starts
 - o client may subscribe/unsubscribe by `matchId`
4. HTTP route mutations invoke `broadcastMatchCreated` / `broadcastCommentary`.
5. WS layer emits events to all clients or scoped match subscribers.

6) Data Model

6.1 `matches`

Represents a sporting fixture and score state.

- `id` (PK)
- `sport`
- `home_team`, `away_team`
- `status` enum: `scheduled` | `live` | `finished`
- `start_time`, `end_time`
- `home_score`, `away_score` (default 0)

- `created_at`

6.2 commentary

Represents timeline events/messages attached to a match.

- `id` (PK)
- `match_id` (FK → `matches.id`)
- `minute`
- `sequence`
- `period`
- `event_type`
- `actor`
- `team`
- `message` (required)
- `metadata` (JSONB)
- `tags` (text[])
- `created_at`

6.3 Relational Notes

- One-to-many: `matches` → `commentary`.
 - Foreign key integrity prevents `commentary` records from referencing missing matches.
-

7) API Surface

7.1 Health / Root

- `GET /`
 - Returns a simple greeting string.

7.2 Matches

- `GET /matches?limit=<n>`
 - Returns latest matches ordered by `createdAt DESC`.
 - Default limit 50, hard cap 100.
- `POST /matches`
 - Validates payload:
 - required: `sport`, `homeTeam`, `awayTeam`, `startTime`, `endTime`
 - optional: `homeScore`, `awayScore`
 - invariant: `endTime > startTime`
 - Computes status using current server time:
 - before start → `scheduled`
 - between start and end → `live`
 - at/after end → `finished`
 - Persists match and broadcasts `match_created` event to WS clients.

7.3 Commentary

Mounted as `app.use('/matches/:id/commentary', commentaryRouter)`.

- `GET /matches/:id/commentary?limit=<n>`
 - Validates path param `id` and optional limit.
 - Returns latest commentary entries for a match by `createdAt DESC`.
 - Default limit 10, hard cap 100.
 - `POST /matches/:id/commentary`
 - Validates `id` plus commentary payload.
 - Inserts commentary row linked to `match_id`.
 - Broadcasts WS `commentary` event to subscribers of that match.
-

8) WebSocket Protocol

Endpoint: `ws://<host>:<port>/ws`

8.1 Server → Client messages

- { "type": "welcome" }
- { "type": "match_created", "data": <match> }
- { "type": "commentary", "data": <commentary> }
- { "type": "subscribed", "matchId": <number> }
- { "type": "unsubscribed", "matchId": <number> }
- { "type": "error", "message": "Invalid JSON" }

8.2 Client → Server messages

- Subscribe:

```
{ "type": "subscribe", "matchId": 123 }
```

- Unsubscribe:

```
{ "type": "unsubscribe", "matchId": 123 }
```

8.3 Subscription Model

- In-memory map: `matchId -> Set<socket>`.
- Each socket stores `subscriptions` set for cleanup on disconnect.
- Broadcast modes:
 - **Global**: new match events to all connected clients.
 - **Scoped**: commentary events to clients subscribed to that match.

Implementation note: in the current `broadcastToMatch` implementation, the `readyState` condition appears inverted (`!== WebSocket.OPEN` before send), which can prevent delivery to active clients; this should be corrected for reliable commentary fan-out.

8.4 Liveness

- Heartbeat every 30 seconds.
- Socket terminated when no pong response is observed.

9) Validation Strategy

Validation is centralized in `src/validation/*` and applied before business logic. Benefits:

- Uniform bad-request response structure.
- Strong guard rails against malformed payloads.
- Coercion for query/path inputs (`z.coerce.number`) where appropriate.

Notable rules:

- Match IDs are positive integers.
- Match list/commentary list limits are positive and capped.
- Match `endTime` must be after `startTime`.
- Commentary `message` is required and non-empty.

10) Security and Abuse Prevention

Arcjet is configured separately for HTTP and WS:

- **Shield**: baseline request hardening.
- **Bot detection**: allows search-engine and preview bots, blocks undesired automation.
- **Rate limiting**:
 - HTTP: 50 requests / 10 seconds / IP.
 - WS upgrade: 5 attempts / 2 seconds / IP.

Failure handling:

- HTTP denied requests return 429 (rate-limited) or 403.
- WS denied upgrades return status line and close socket.
- Arcjet errors return 503 (HTTP) or 500 (WS upgrade).

11) Configuration and Environment

Required environment variables:

- `DATABASE_URL` — PostgreSQL connection string.
- `ARCJET_KEY` — Arcjet API key (required by current implementation).

Optional:

- `ARCJET_MODE` — set to `DRY_RUN` for non-blocking policy evaluation; defaults to `LIVE`.
- `PORT` — defaults to `8000`.
- `HOST` — defaults to `0.0.0.0`.

Operational note: startup currently throws if `ARCJET_KEY` is missing.

12) Database Migrations and Schema Management

Drizzle workflow (from `package.json`):

- `npm run db:generate` — generate migration SQL from schema.
- `npm run db:migrate` — apply migrations to DB.
- `npm run db:studio` — open Drizzle Studio.

`drizzle.config.js` points Drizzle Kit to:

- schema source: `src/db/schema.js`
- output folder: `drizzle/`
- dialect: PostgreSQL

13) Operational Characteristics

- Clear separation of concerns (routing, validation, data, WS, security).
- Real-time updates integrated without third-party brokers.
- Strong input validation posture via Zod.
- Schema-driven persistence and migration process.

14) Component-by-Component Deep Dive

`src/index.js`

- Bootstraps Express + HTTP server.
- Registers middleware and routers.
- Instantiates WS server and stores broadcast functions in `app.locals` so HTTP handlers can trigger real-time messages.

`src/arcjet.js`

- Constructs two Arcjet clients:
 - `httpArcjet` for Express middleware checks.
 - `wsArcjet` for upgrade request protection.
- Exposes `securityMiddleware()` used globally on HTTP routes.

`src/db/db.js`

- Loads environment config.
- Creates a shared `pg.Pool` and wraps it with Drizzle.

`src/db/schema.js`

- Defines schema tables and enum in code-first style.
- Serves as the canonical source used by Drizzle Kit to generate migrations.

`src/routes/matches.js`

- Implements list/create endpoints for matches.
- Enforces payload/query rules with Zod.
- Derives status from schedule windows.
- Broadcasts newly created matches over WS.

`src/routes/commentary.js`

- Implements list/create commentary under a match context.
- Validates both route params and request body.

- Broadcasts inserted commentary to WS subscribers of target match.

`src/validation/matches.js / src/validation/commentary.js`

- Central contract layer for incoming API traffic.
- Encodes type checks, coercions, and business invariants.

`src/utils/match.status.js`

- Encapsulates schedule-to-status derivation logic.
- Contains helper to synchronize persisted status if stale.

`src/ws/server.js`

- Manages WS lifecycle:
 - upgrade gate
 - connect/disconnect
 - heartbeat
 - subscription bookkeeping
- Exposes broadcast methods consumed by HTTP route layer.

15) Suggested Quick Start (Developer)

1. Create `.env` with `DATABASE_URL`, `ARCJET_KEY`, and optionally `ARCJET_MODE`, `PORT`, `HOST`.
2. Install dependencies: `npm install`.
3. Generate/apply schema: `npm run db:generate && npm run db:migrate`.
4. Run server: `npm run dev` (or `npm start`).
5. Open WS client to `/ws` and subscribe to match IDs for live commentary stream.